

# Elements of DeFi

<https://web3.princeton.edu/elements-of-defi/>

**Professor** Pramod Viswanath

Princeton University

# **Lecture 3: Smart Contracts and Pricing**

# Last time

- Blockchains to track transfer of tokens, maintain a ledger

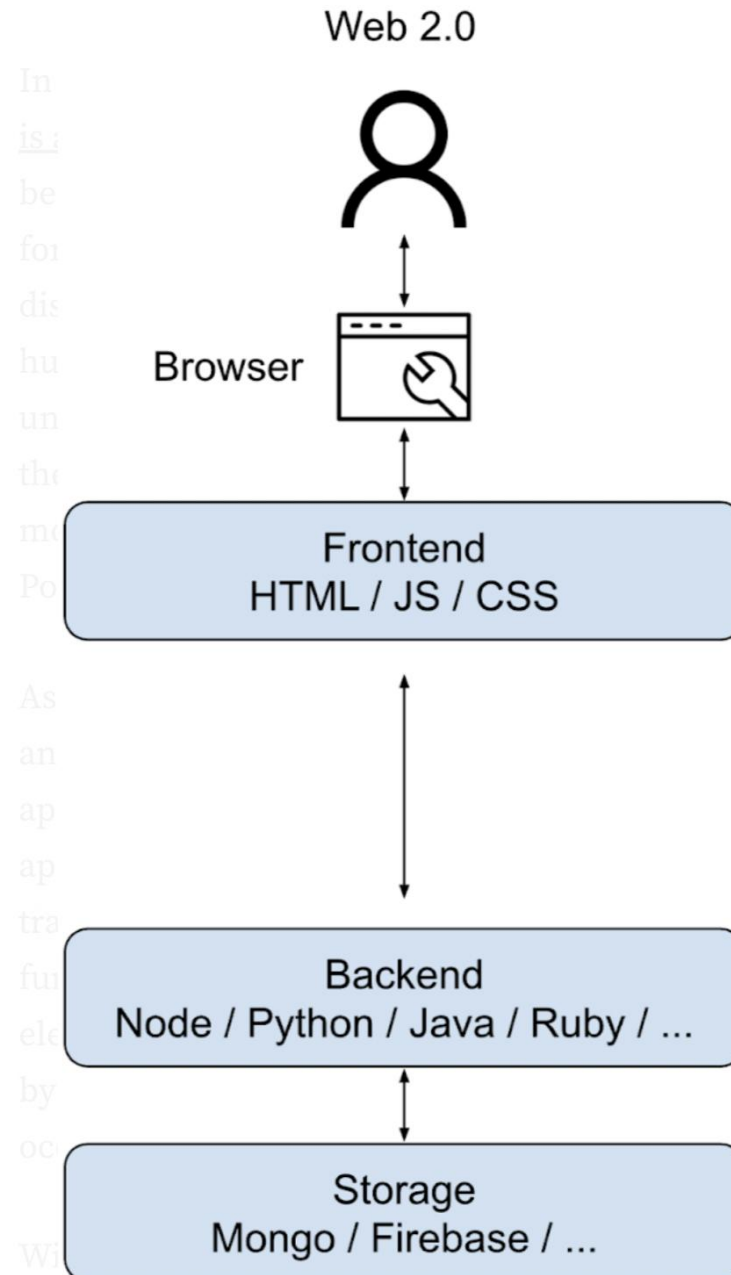
# This time

- Blockchains to run programs, store information, transfer tokens,
- Tool used : “smart contracts”
  - computer programs that use the ledger entries as variables and memory
- A “transaction” is now generalized into a “function call”
  - smart contract
- Ledger update is managed via a virtual machine
  - Generalize the simple ledger to a “state machine”, a run-time environment for smart contract execution
  - Ethereum Virtual Machine
- Auctions for selling transaction slots

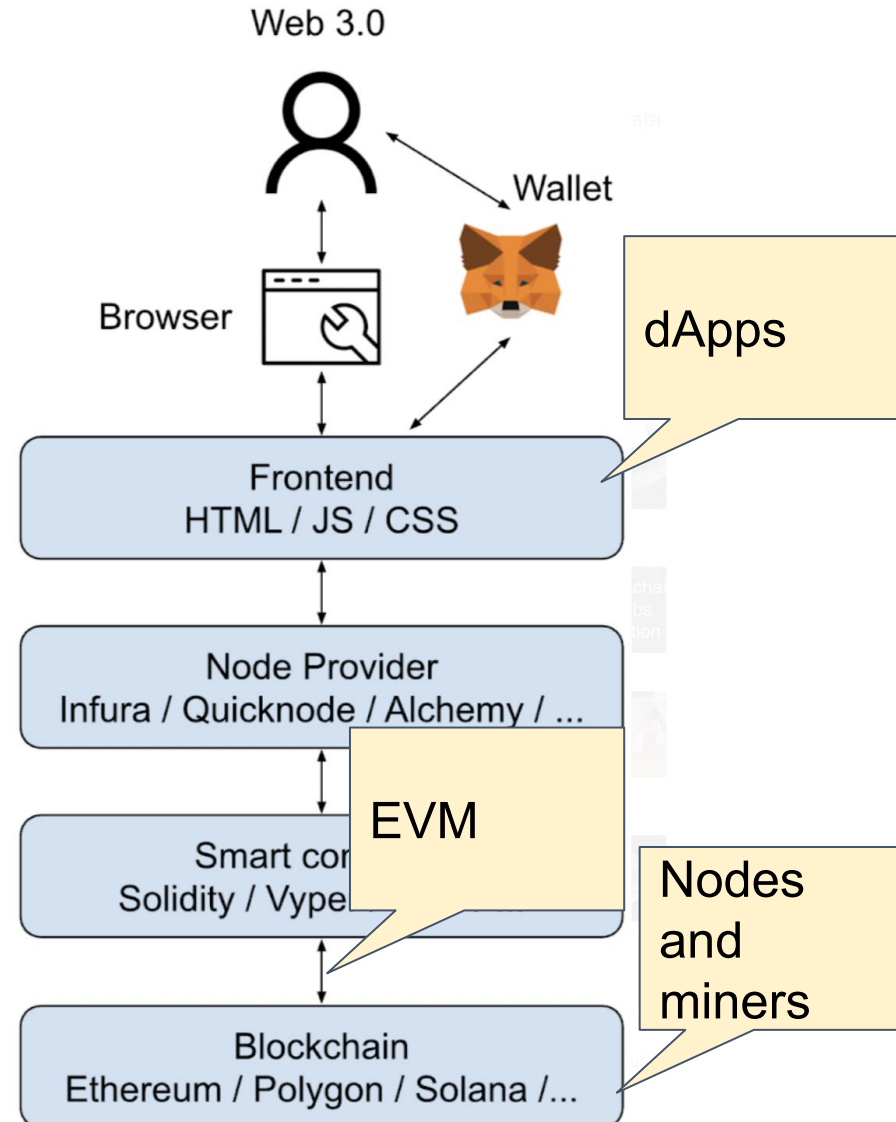
# Outline

- What are smart contracts?
- How is smart contract programming different?
  - web 3.0 development
  - dApp programming
- EVM
  - opcodes, gas fees
- Solidity
- Pricing of transactions
- Today's lab :
  - Fungible tokens: ERC20

# Web 2.0 : a system view




# Web 3.0 : a system view



# What is a Smart Contract?

- First proposed in 1990s, “digital form of promises”
- Not necessarily related to a contract
- Lifecycle:
  - Deployed by a transaction
  - Establish initial states, immutable once deployed
  - Store states and execute computations
- Languages: Solidity, Rust, JavaScript...





# Where is the program stored?

1. Writing a program - using high-level language such as Solidity, Vyper, ...
2. Program is compiled - into low-level language that the blockchain state machine can understand
3. Program is deployed - binary is stored in the state of the blockchain

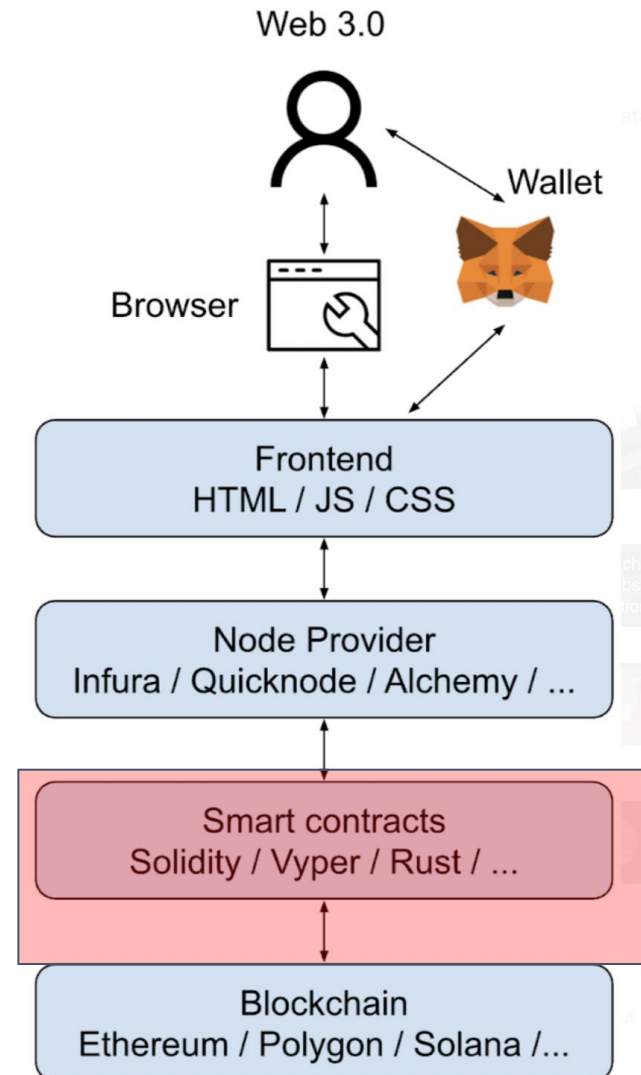
# What kind of a programming language should be used?

- Should be similar to existing languages in syntax
- Should have additional functionality to access blockchain state
- **Every execution should yield the same result**
  - so execution output can be verified by everyone else
  - rand() not allowed
  - function calls to outside (e.g., the internet) are not allowed
- Cost of execution (gas fee) should be readily calculated

# Design goals as a developer

- Correctness –
  - especially important because of Byzantine actors - program forced to go into a “rare” corner case
  - stakes are high! - cannot change code once deployed
  - tokens handled have value - incentive for bad actors
- Efficiency –
  - one pays fees for each execution of the smart contract and this adds up on each execution

# This lecture

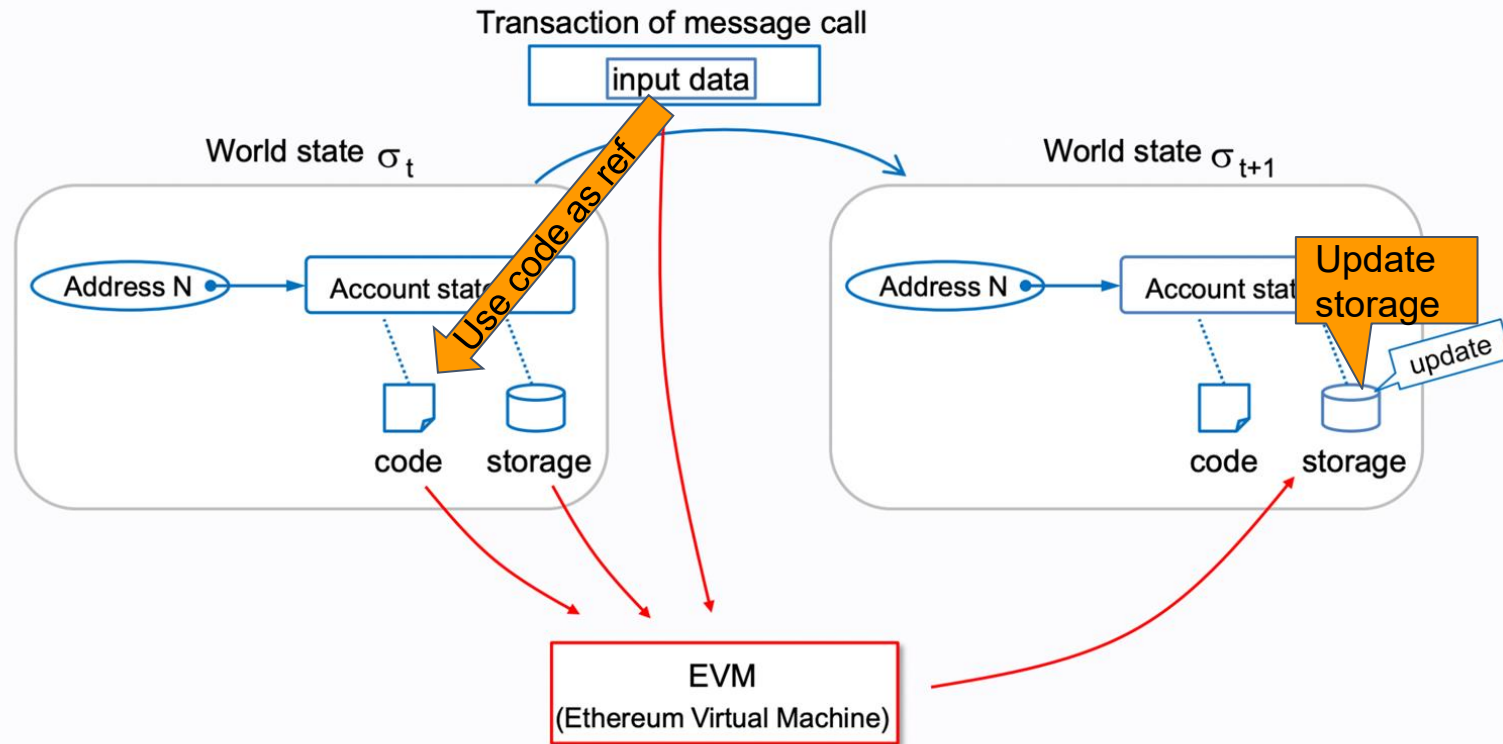


# EVM

- The Ethereum Virtual Machine is the **distributed** execution environment (“state machine”) running on the Ethereum blockchain
- Each block on Ethereum changes the state of EVM
- Every Ethereum user sees the same canonical EVM state at any given block

# EVM: state updates

- State of EVM changed via transactions :



EVM code is executed on Ethereum Virtual Machine (EVM).

# EVM: transactions, opcodes

- Two types of transactions :
  - Resulting in function calls
  - Resulting in contract creation
- Every transaction is decomposed into a sequence of OPCODEs
  - e.g. ADD, SUB, JUMP, LOAD, ...
  - fixed number (256) of opcodes
- Every OPCODE consumes a fixed amount of gas
  - total gas of a transaction is the sum of gas of constituent opcodes
  - gas to eth is a variable

# EVM: gas

Table 1. EVM opcodes and gas cost

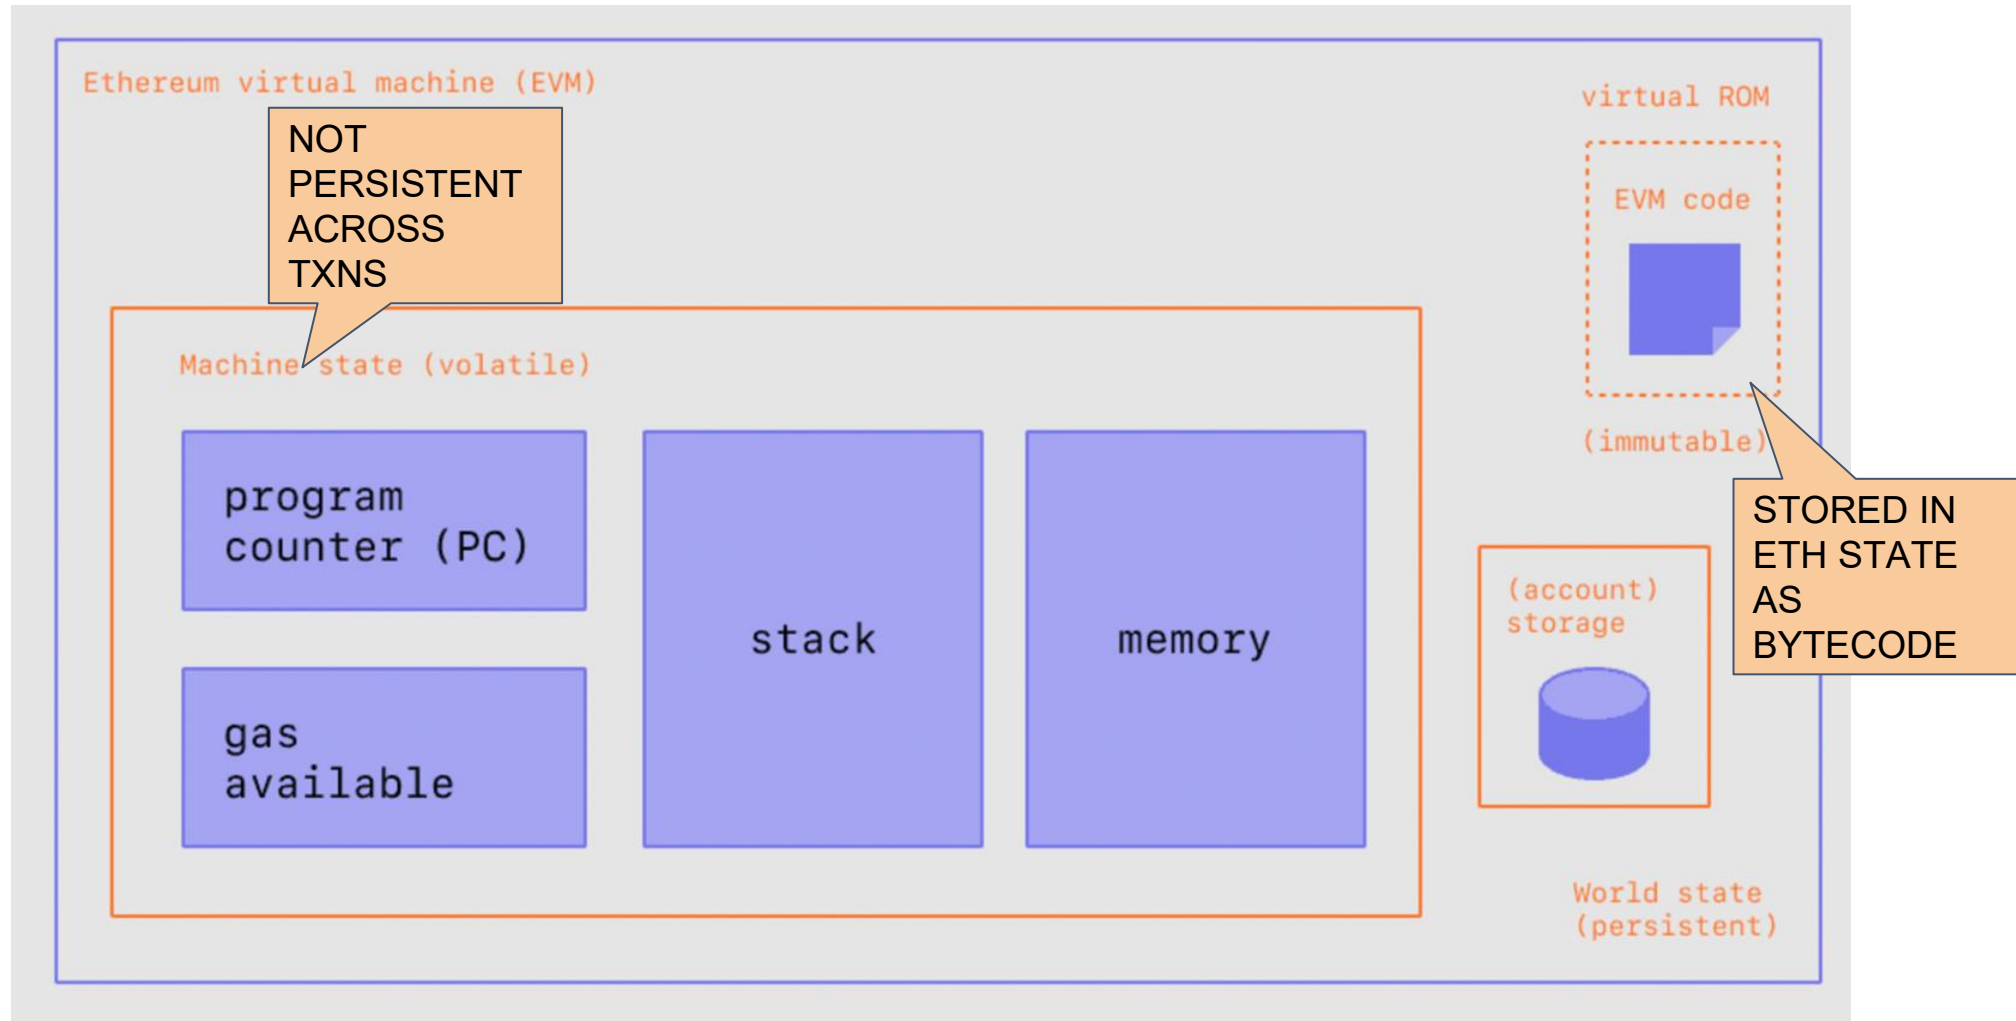
Opcode	Name	Description	Extra info	Gas
0x00	STOP	Halts execution	-	0
0x01	ADD	Addition operation	-	3
0x02	MUL	Multiplication operation	-	5
0x03	SUB	Subtraction operation	-	3
0x04	DIV	Integer division operation	-	5
0x05	SDIV	Signed integer division operation (truncated)	-	5
0x06	MOD	Modulo remainder operation	-	5
0x07	SMOD	Signed modulo remainder operation	-	5
0x08	ADDMOD	Modulo addition operation	-	8
0x09	MULMOD	Modulo multiplication	-	8



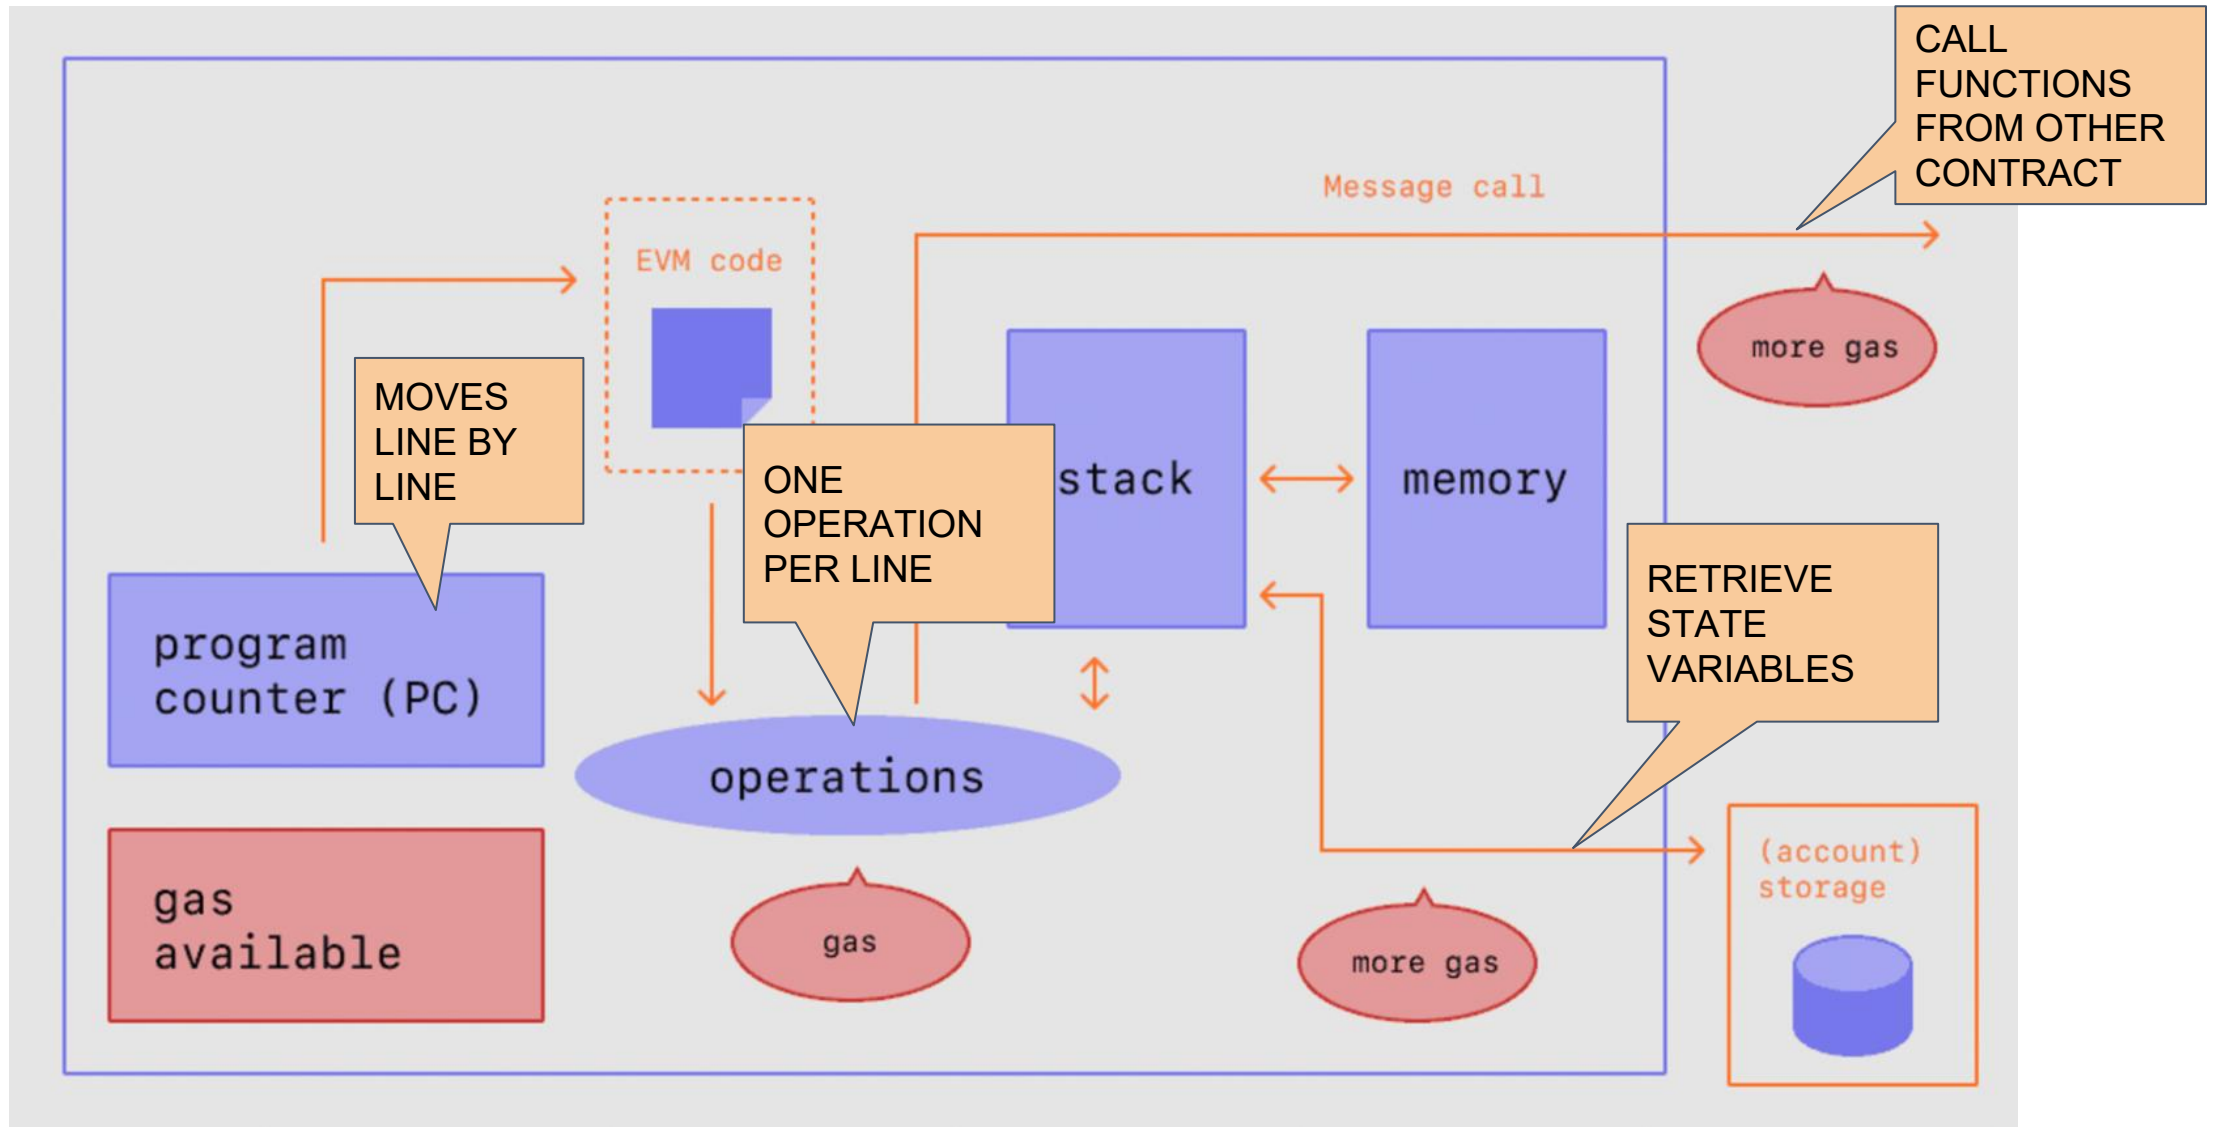
# EVM: Data Store

- Storage
  - Written in the blockchain, stored permanently
  - Expensive, some gas is refunded when storage is deleted
- Memory
  - A byte array with slot sizes of 32 bytes
  - Stored during function execution
  - Cheap, but the costs per operation scales quadratically
  - Does not persist across txns
- Stack
  - Only 16 stack variables are accessible
  - Cheapest, manipulated by inline assembly

# Tracing a transaction through the data store



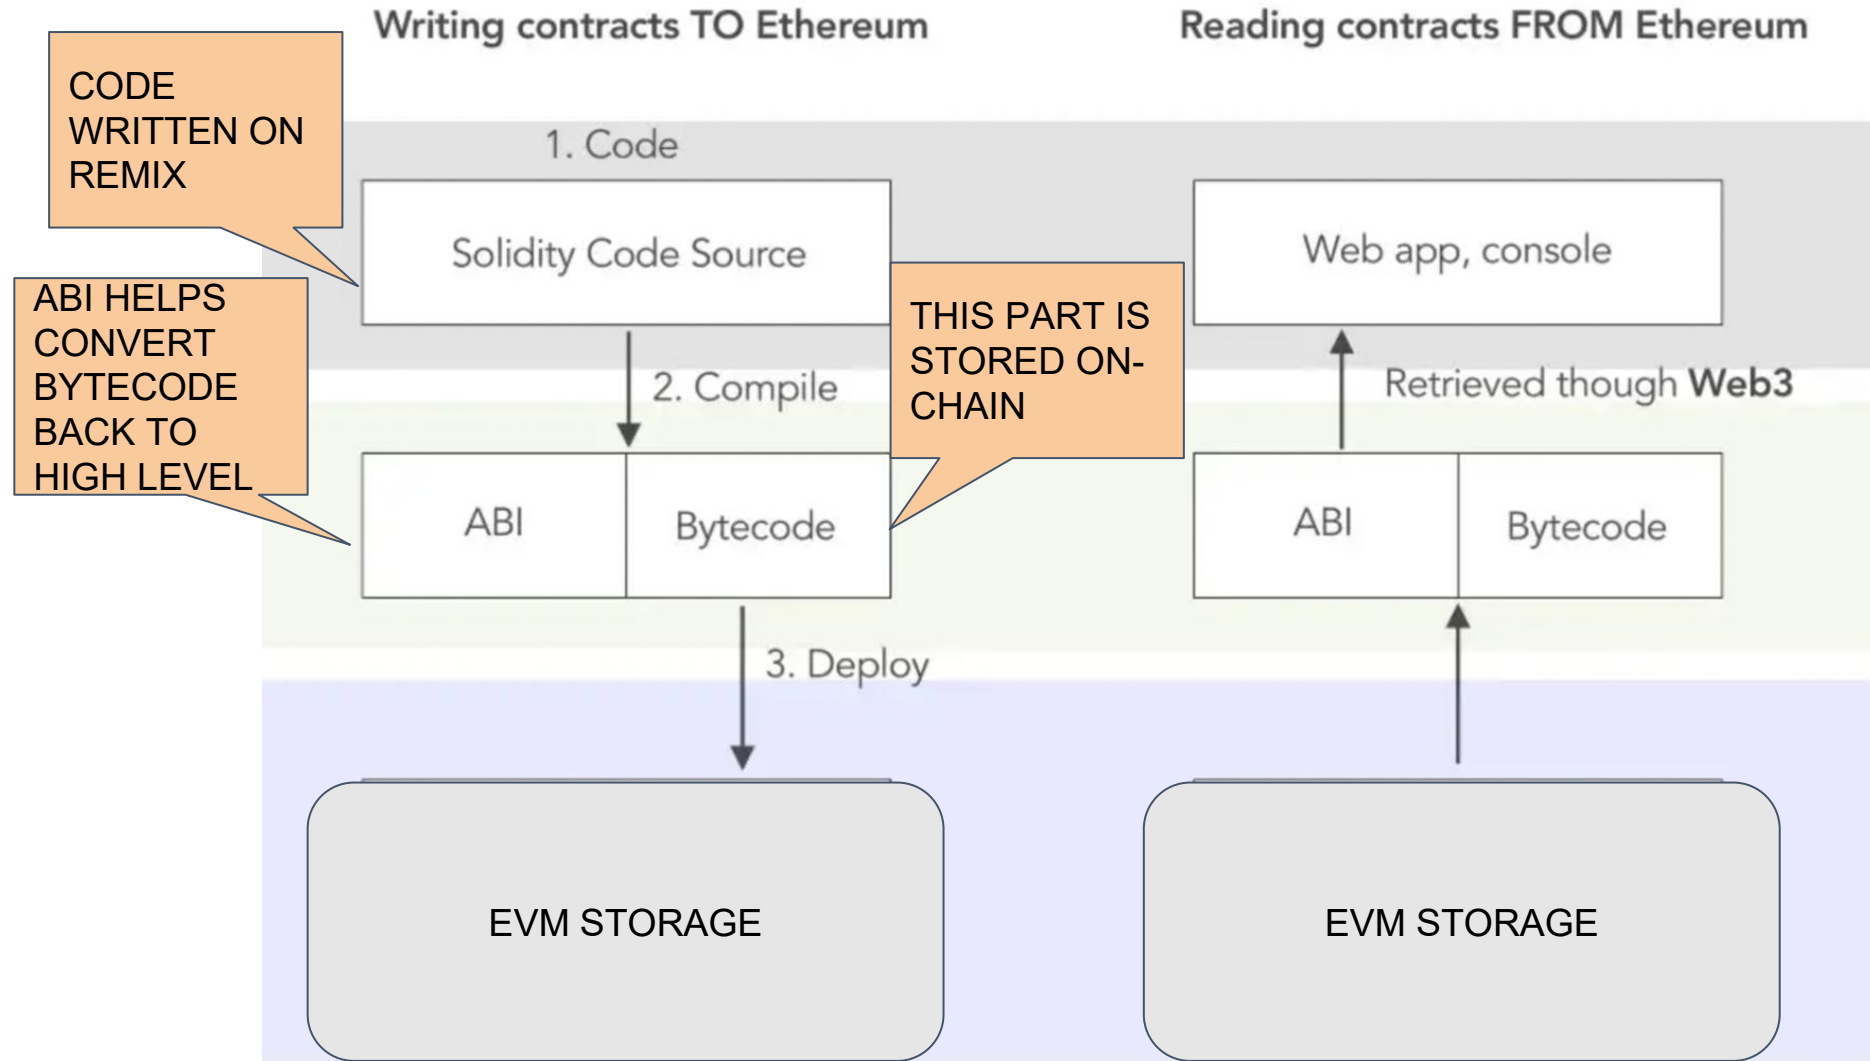
# Tracing a transaction through the data store



# Solidity

- Object-oriented, statically-typed
- Designed for Ethereum
  - Also used by Binance Smart Chain, Avalanche, XinFin...
- Turing-complete
- Popular IDE – remix (we use this in our lab, assignments)

# Solidity : Interaction with EVM



# Solidity Bytecode Opcode

```
60 00          = PUSH1 0x00
35             = CALLDATALOAD
60 e0          = PUSH1 0xe0
1c            = SHR
80            = DUP1
63 2e64cec1    = PUSH4 0x2e64cec1
14            = EQ
61 003b        = PUSH2 0x003b
57            = JUMPI
80            = DUP1
63 6057361d    = PUSH4 0x6057361d
14            = EQ
61 0059        = PUSH2 0x0059
57            = JUMPI
```

CONTRACT STORED  
AS BYTECODE

INTERPRETED BY EVM

AT EXECUTION,  
BYTECODE CONVERTED  
TO OPCODES FOR THE  
EVM

```
26   return number;
27   }
28 }
```

# Problem

- EVM Compute + Storage + Memory are scarce resources

Therefore, **gas** is a scarce resource

- How should txn submitters bid for gas?
- Some kind of an auction? How should the auction be settled?
- We look at this question from the perspective of **miner incentives**


# Design 1 : a first-price auction

- (Implemented in Bitcoin and Ethereum basic)
  1. Every txn submits a bid
  2. Miner includes the N highest bids
  3. Everyone pays fees equal to their bids
  4. All of the money goes to the miner
- Called a “first-price auction”



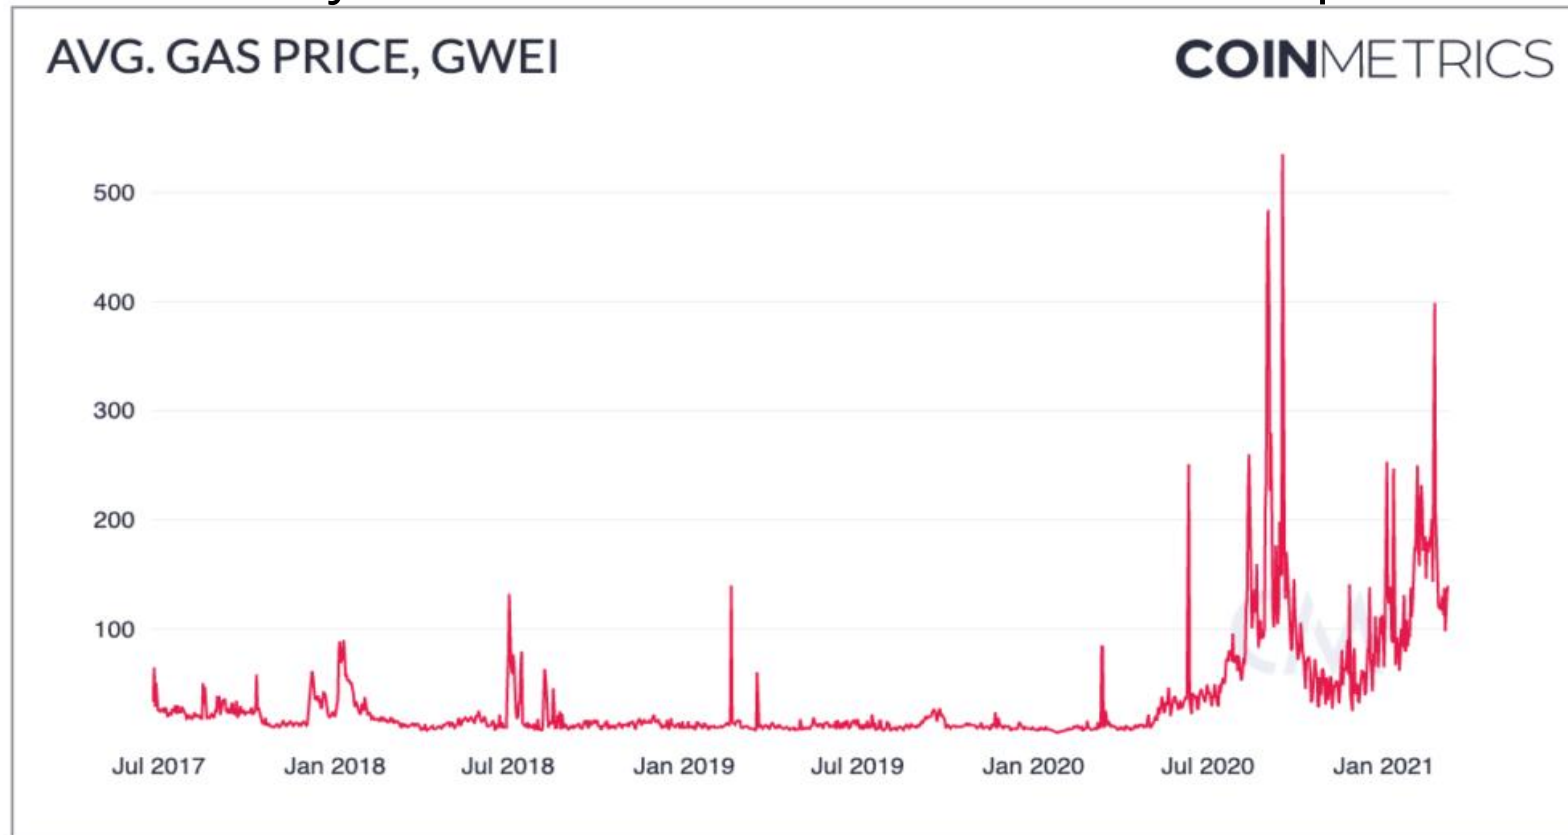
# First-price auctions – what's **wrong**?

- **Lack of Incentive compatibility**
  - Not clear how much to bid
  - Overpayment of fees – user ends up thinking of what other people might bid
  - Need to estimate prices – complex and still results in overpayment




# First-price auctions – what's **wrong**?

- **Bidding up of prices – does not match actual txn cost**
  - Txn costs skyrocket to unreasonable amounts in practice




# First-price auctions – what's **wrong**?

- **Blockchain instability**
  - Even after block has been mined, other miners attempt to undercut
  - Dominant mining strategy : deviate from protocol
  - This makes a “51% attack” achievable at lesser hash power




# Design 2 : a second-price auction

1. Every txn submits a bid
2. Miner includes the N highest bids
3. Everyone pay fees equal to N+1-highest bid
4. All of the money goes to the miner




# Second-price auction – what's **right**?

- **Incentive compatibility**
  - Bidding is easy now
  - Simply bid how much value you have in mind for the txn



# Second-price auction – what's **wrong**?

- Miner can increase fee charged - by inserting txns
  - Miner can insert their own txns
  - Miner bids just below Nth bid OR
  - Miner introduces  $N/2$  txns below  $N/2$ th bid ...



# Second-price auction – what's **wrong**?

- **Miner played by the rules – and earned extra money**
  - Such behavior makes user experience worse and markets inefficient
  - Profit obtained is our first encounter with “MEV”
  - MEV – Miner Extractable Value


# Design 3 : Monopolistic Auction

1. Every txn submits a bid
2. Calculate  $N^* = \operatorname{argmax} (N\text{th highest bid}) \times N$
3. Miner includes the top  $N^*$  highest bids
4. Everyone pays fees equal to  $N^*$  th-highest bid
5. All of the money goes to the miner




# Monopolistic Auction – what's **right**?

- **Incentive compatibility remains**
  - Bidding : no incentive to bid higher or lower
  - Caveat :  $N^{\text{th}}$  person incentivized to bid slightly lower, but not a big difference irl




# Monopolistic Auction – what's **right**?

- **Miner conducts auction honestly**
  - Miner gains no profit by inserting its own transactions
  - Proof is non-trivial, but intuition is – if miner tries to drive up price, then  $N^*$  decreases  $\rightarrow$  revenue collected stays the same



# Monopolistic Auction – what's **wrong**?

- **Off-chain collusion**
  - Miner gains profit by eliciting bids off-chain first
  - Miner can make an offer that is beneficial to everyone!
  - How? Miner gets \$14 instead of \$12




# Design 4 : EIP 1559

1. Price  $p$  fixed by the protocol “base fee” + can include an optional tip
2. Miner picks at most  $N$  txns
3. Amount  $p$  from each user is **burnt**
4. All tips collected + a fixed block reward go to miner


# EIP-1559 – what's **right**?

- **Incentive compatibility remains**
  - If  $p$  is chosen so that  $< N$  users have utility  $> p$ , then it is incentive compatible
  - Bid only if utility  $> p$ , tip = 0



# EIP-1559 – what's **right**?

- **Miners are no longer incentivized to cheat or collude**
  - If  $p$  is chosen so that  $< N$  users have utility  $> p$ , then everyone charged a fixed price and miner gets a fixed block reward
  - Colluding off-chain would not be beneficial for users



# EIP-1559 – what's **wrong**?

- What happens when the value of  $p$  is off? i.e.  $> N$  users have utility  $> p$ 
  - Get back the first price auction
  - This is still resistant to miner cheating or collusion
  - But no longer incentive compatible
- Protocol needs to choose base fee  $p$  carefully

# EIP-1559 – updating base fee

1. Define two constants :
  1. *maximum block size* = **N**
  2. *target block size* = **N/2**
2. Update p according to previous block size
  1. Increase p when previous block size > target
  2. Decrease p when previous block size < target
3. Uses the following rule to do that :

$$p_{new} = p_{old} \left( 1 + \frac{1}{8} \frac{B_{prev} - B_{target}}{B_{target}} \right)$$



# EIP-1559 – updating base fee

- Current protocol update rule still a bit ad-hoc
- Open questions :
  - Best way to update the base fee  $p$ ?
  - Better mechanisms to collect fees?

# EIP-1559 – what's **wrong**?

- Usually the tip revenue is minor – but all the rest of the revenue gets burnt
- Miner not incentivized to process transactions
- Miner goes off-chain to “extort” – announces that will only include txn if paid some amount xyz
- Turns out – **no auction exists that is incentive compatible, credible, collusion-proof, and extortion-proof**
- **Open problem** – how do you effectively trade-off between these four properties?

# EIP-1559 – what's **wrong**?

- EVM consumes distinct kinds of resources : Compute, Storage, Memory
- All resources are quantified in the same terms : gas fees
- **Problem?**
- Suppose block has many txns that are CPU-intensive
- Competition for CPU drives up gas fees
- Drives up costs of doing other non-CPU txns as well – txns consuming normal amounts of bandwidth or memory or storage

# Multi-dimensional pricing

- Keep updating a *vector* of prices  $\mathbf{p}$
- Number of dimensions = Number of distinct resources
- Use vector form of update equation :
$$p_{new} = p_{old} \left( 1 + \frac{1}{8} \frac{B_{prev} - B_{target}}{B_{target}} \right)$$
- Here  $B_{prev}$  and  $B_{target}$  are also fixed vectors set by protocol and evolved slowly
- Turns out, the update equation above is same as doing gradient descent to maximize blockchain user welfare – **is essentially optimal way to set prices**

# Summary

- Smart contract introduction
- EVM + Solidity + Gas
- Pricing of transactions on the EVM

# Next Lecture

Meet our first **element** - Exchanges